

Host the tool & enable “Suggest Targeting”

Turn campaign plans into fully-built FreeWheel orders — and give the team AI-assisted targeting from one shared URL.

15 / ~79

REGIONS / BRANDS

389

TESTS PASSING

Live

FREEWHEEL PROD (DRAFTS)

1 VM

INFRA ASK

The decision: stand up one small internal host so our 5-person team uses the Campaign Plan form + “Suggest Targeting” from a single URL — instead of everyone running Python on their laptop.

WHAT WE’RE HOSTING — TWO SEPARABLE PIECES

The immediate ask

🌟 **Suggest Targeting helper**

A ~130-line stdlib Python web app. Serves the form + answers one /suggest endpoint. Proposes grounded targeting (comp shows, Pluto channels, genres) from our own past campaigns.

- **No database, no FreeWheel credentials.**
- Behind VPN/SSO, one port (8770).

Later / optional

Order-building CLI & automation

Builds & pushes FreeWheel orders (drafts only). Needs FreeWheel credentials. Can stay operator-side for now; moves server-side once Salesforce + MRM creds land.

- Creates **NOT_BOOKED drafts only** — a human activates.
- Daily data-refresh job already scripted.

SUGGEST RUNS IN TWO MODES — THIS IS THE SECURITY CRUX

Start here · no approval needed

Historicals-only (default)

Suggestions come only from our own past-campaign data. **Zero external calls — nothing leaves our environment.** Works today.

Add later

AI-enhanced (optional)

One **server-side** ANTHROPIC_API_KEY adds an AI layer for brand-new titles. Key never touches a browser. Only this needs the API-key sign-off.

WHAT WE NEED FROM ENGINEERING

Host	One always-on Linux VM/container — 1 vCPU, 512 MB–1 GB RAM . No GPU, no database.
Runtime	Python 3.10+. Pull the repo, run one command.

Network	Reachable by 5 users behind VPN/SSO . Not public. One TCP port.
Outbound	Historicals-only: none . AI layer: HTTPS to Anthropic only. CLI push: HTTPS to FreeWheel.
Process	systemd/container for auto-restart + a daily refresh - form job.
Secrets	Env/secret store for the optional Anthropic key and (if CLI runs here) FreeWheel creds.

RECOMMENDED ROLLOUT

1

Now
Host Suggest in **historicals-only** mode behind VPN. Immediate value, no blockers, no external traffic.

2

Next
Add the **Anthropic key** once approved — same host, just a restart. Turns on the AI layer.

3

Later
Move order-building / Case automation server-side once **Salesforce + FreeWheel MRM** creds are provisioned.

Bottom line: one small internal VM + one command unblocks the whole team today, with nothing leaving our environment. The optional AI layer is a single server-side key we can switch on later — no rework.

OPEN DEPENDENCIES TO NAME IN THE ROOM

- **Salesforce** field/credential setup — for full Case → drafts automation (building from the form/CLI doesn't need it).
- **FreeWheel MRM client-credentials** — until then the IO Brand is set by hand on push; everything else works.
- **API-key sign-off** — the only gate on the Suggest AI layer.