

Automated order building + Suggested Targeting

The build is done and running against FreeWheel production. Putting the last mile on the roadmap — hosting it + turning on Suggested Targeting — converts it into faster, more accurate order building for the whole team.

–75%

Error rate — pre-built, rule-complete, validated before push

–60%

Time to generate & build an order

15 / ~79

Regions / brands already encoded

389

Automated tests guarding behavior

THE PROBLEM TODAY

Promo campaigns are trafficked **by hand** — dozens of placements per order across tiers, durations, formats, regions and platforms, each with precise targeting, ad units, geo, flight timing and exclusions. Coordinators rely on **institutional knowledge, common sense, and scattered reference docs**. That's slow, error-prone, inconsistent, and hard to scale or onboard.

WHAT CHANGES

TODAY — MANUAL

- Dozens of settings keyed by hand per order
- Targeting from judgment + manual lookups
- Errors surface after campaigns go live
- Output varies by who built it
- Rules live in people's heads

WITH THE TOOL

- Rule-complete draft generated in seconds
- Targeting grounded in real inventory + history
- Malformed orders blocked **before** push
- Same rules every time — consistent & auditable
- Best practice encoded once, in the system

WHY IT MOVES THE NEEDLE

**–75%
errors**

Rule-complete by construction

The settings people mis-key today — site groups, ad units, exclusions, rating VGs, flight dates — are applied consistently and validated before push. Fewer live-campaign incidents, make-goods, and re-trafficking cycles.

–60% time

Review, don't rebuild

Significant manual setup becomes a reviewed draft in seconds. Coordinators shift from building placements to QA — absorbing more volume without adding headcount.

Accuracy

Consistent & auditable

Every order follows the same rule set, mirrored from the live reference IOs. Deterministic and

**Smarter
targeting**

Real data, not guesswork

reviewable — and the tool never books on its own; a human activates every draft.

Suggested Targeting proposes comp shows, Pluto channels, genres & categories — matched to our real FreeWheel inventory and what similar past campaigns actually ran. Huge for new titles and onboarding.

Illustrative: at ~1,000 orders/yr at ~45 min each, a 60% cut saves **~450 hours/year** — before counting the error-rate savings, which avoid the costliest failures (live mis-targeting). Swap in our real volumes to finalize.

SPOTLIGHT — THE SUGGESTED TARGETING TOOL

Data-based targeting direction — strongest when other resources fall short

Targeting is where campaigns win or waste impressions, and today it's built from judgment and scattered references. Suggested Targeting gives every coordinator a **strategic, intentional, data-grounded starting point** instead of guesswork.

- **More strategic & intentional** — proposes comp shows, Pluto channels, genres & categories *with a rationale*, ranked by agreement across signals, not gut feel.
- **Grounded in real data** — every suggestion is matched to our live FreeWheel inventory and to what *similar past campaigns in that region actually ran*.
- **Most valuable exactly when references are thin** — new titles with no history, unfamiliar regions or genres, or tight timelines, where a coordinator would otherwise be guessing.
- **Levels up the whole team** — encodes best-practice targeting so a newer coordinator builds like our most experienced one.

 **Requires the external AI component to work as intended.** The AI layer is what reasons about an unfamiliar title, layers in the brief, and turns our real data into a usable recommendation — it's the piece that makes the guidance genuinely strategic. It runs on **one server-side key**; no data sits in anyone's browser.

THE ASK — A SMALL, WELL-SCOPED LAST MILE

1

Host it

One small internal VM (behind VPN) serves the form + Suggested Targeting to all 5 of us.

2

Enable Suggest

Turn on Suggested Targeting — the external AI component powers the strategic guidance (one server-side key; a historicals-only mode is a limited fallback).

3

Own the refresh

A daily data-refresh job (already scripted) keeps dropdowns & suggestions current.